

Lab 07: Support Vector Machines

Today we'll look at the SVM maximum margin classification problem and how we can implement the optimization in Python.

We'll use the cvxopt quadratic programming optimizer in Python.

Later in the lectures, we'll see that more specialized algorithms such as Sequential Minimal Optimization implemented by the machine learning libraries are more effective for large SVM problems.

Linearly separable case: direct solution using quadratic programming

If we assume that the data are linearly separable, we can use the following setup for the optimization:

- The data are pairs $(\mathbf{x}^{(i)}, y^{(i)})$ with $\mathbf{x}^{(i)} \in \mathbb{R}^n$ and $y^{(i)} \in \{-1, 1\}$.
- The hypothesis is

$$h_{\mathbf{w},b}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}^\top \mathbf{x} + b > 0 \\ -1 & \text{otherwise} \end{cases}$$

- The objective function is

$$\mathbf{w}^*, b^* = \operatorname{argmax}_{\mathbf{w}, b} \gamma,$$

where γ is the minimum geometric margin for the training data:

$$\gamma = \min_i \gamma^{(i)}$$

and $\gamma^{(i)}$ is the geometric margin for training example i , i.e., the signed distance of $\mathbf{x}^{(i)}$ from the decision boundary, with positive distances indicating that the point is on the correct side of the boundary and negative distances indicating that the point is on the incorrect side of the boundary:

$$\gamma^{(i)} = y^{(i)} \left(\left(\frac{\mathbf{w}}{\|\mathbf{w}\|} \right)^\top \mathbf{x}^{(i)} + \frac{b}{\|\mathbf{w}\|} \right).$$

- To find the optimal \mathbf{w}, b according to the objective function above, we can solve the constrained optimization problem

$$\begin{array}{ll} \min_{\mathbf{w}, b} & \gamma \\ \text{subject to} & y^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} + b) \geq 1, \quad i = 1, \dots, m \end{array}$$

```
In [41]: # in case of there is not cvxopt
!pip install cvxopt
```

```
Requirement already satisfied: cvxopt in /home/alisa/anaconda3/lib/python3.8/site-packages (1.2.7)
```

```
In [42]: import numpy as np
import matplotlib.pyplot as plt
import cvxopt
```

The example data

```
In [43]: Xf = np.matrix([[ 164.939, 163.431, 157.554, 152.785, 156.385, 15
9.242, 156.281, 164.411, 157.308, 159.579 ],
                        [ 56.927, 48.945, 45.678, 45.969, 40.896, 4
6.848, 42.225, 42.380, 42.150, 49.739 ]]).T;
Xm = np.matrix([[ 168.524, 171.597, 179.469, 176.063, 180.939, 17
7.011, 183.284, 180.549, 176.502, 185.392 ],
                [ 64.353, 61.793, 74.552, 69.851, 74.730, 7
5.871, 79.170, 79.753, 64.923, 73.665 ]]).T;
X = np.concatenate([Xf, Xm],0);
y = np.concatenate([-np.matrix(np.ones([10,1])),np.matrix(np.ones
([10,1]))]);
```

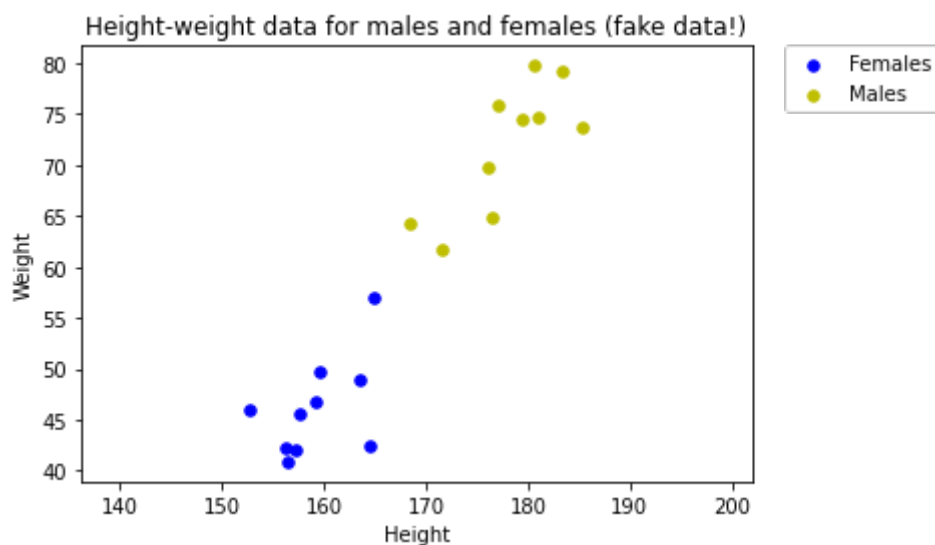
```
In [44]: print(X.shape)
          print(y.shape)
```

```
(20, 2)
(20, 1)
```

```
In [45]: # Plot the data
def plot_mf(Xf,Xm):
    axes = plt.axes()
    females_series = plt.scatter(np.array(Xf[:,0]), np.array(Xf[:,1]), s=30, c='b', marker='o', label='Females')
    males_series = plt.scatter(np.array(Xm[:,0]), np.array(Xm[:,1]), s=30, c='y', marker='o', label='Males')
    axes.set_aspect('equal', 'datalim')
    plt.xlabel('Height')
    plt.ylabel('Weight')
    plt.title('Height-weight data for males and females (fake data!)')
    plt.legend(handles=[females_series, males_series], bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
    return axes

def plot_w(w,b):
    ylim = plt.axes().get_ylim()
    xlim = plt.axes().get_xlim()
    p1 = (xlim[0], - (w[0,0] * xlim[0] + b) / w[1,0])
    p2 = (xlim[1], - (w[0,0] * xlim[1] + b) / w[1,0])
    plt.plot((p1[0],p2[0]), (p1[1],p2[1]), 'r-')

plot_mf(Xf,Xm)
plt.show()
```



Exercise 1 (in lab): linearly separable data (Total 25 points)

Take the example data and SVM optimization code using cvxopt from the exercise in lecture. Verify that you can find the decision boundary for such "easy" cases. Show your results in your lab report.

Exercise 1.1 Create SVM function using cvopt (5 points)

```
In [46]: def cvxopt_solve_qp(Q, c, A=None, B=None, E=None, d=None):
# Fill your code value in 'None'
# Some 'None' can be avoided.
Q_new = None
args = [None, None]
if A is not None:
    args.extend([None, None])
    if E is not None:
        args.extend([None, None])
sol = None
if sol is not None and 'optimal' not in sol['status']:
    return None
x = None
### BEGIN SOLUTION
Q_new = .5 * (Q + Q.T) # make sure Q is symmetric
args = [cvxopt.matrix(Q_new), cvxopt.matrix(c)]
if A is not None:
    args.extend([cvxopt.matrix(A), cvxopt.matrix(B)])
    if E is not None:
        args.extend([cvxopt.matrix(E), cvxopt.matrix(d)])
sol = cvxopt.solvers.qp(*args)
if 'optimal' not in sol['status']:
    return None
x = np.array(sol['x']).reshape((Q.shape[1],))
### END SOLUTION
return x
```

```
In [47]: # Test function: Do not remove
Q_test = np.array([[1, 0],[0, 0]])
c_test = np.zeros([2])
A_test = np.array([[15., 1.],[14., 1.], [13., 1], [-5.,-1.],[-6.,-1.],[-8,-1]])
B_test = -np.ones([6])
try:
    x_test = cvxopt_solve_qp(Q_test, c_test, A_test, B_test)
    print('x_test:', x_test)
except:
    assert False, "cvxopt_solve_qp is incorrect"

print("success!")
# End Test function
```

	pcost	dcost	gap	pres	dres
0:	2.8800e-02	1.0464e+00	1e+01	1e+00	3e+01
1:	1.8859e-01	-8.5852e-01	1e+00	2e-01	3e+00
2:	1.6523e-01	3.7852e-02	1e-01	4e-16	4e-16
3:	8.2300e-02	7.2465e-02	1e-02	6e-16	3e-15
4:	8.0043e-02	7.9915e-02	1e-04	7e-16	2e-16
5:	8.0000e-02	7.9999e-02	1e-06	2e-16	1e-15
6:	8.0000e-02	8.0000e-02	1e-08	8e-16	1e-15

Optimal solution found.
x_test: [-0.40000001 4.20000012]
success!

Expect result (or look-alike): \ pcost dcost gap pres dres \ 0: 2.8800e-02 1.0464e+00 1e+01 1e+00 3e+01 \ 1: 1.8859e-01 -8.5852e-01 1e+00 2e-01 3e+00 \ 2: 1.6523e-01 3.7852e-02 1e-01 4e-16 4e-16 \ 3: 8.2300e-02 7.2465e-02 1e-02 6e-16 3e-15 \ 4: 8.0043e-02 7.9915e-02 1e-04 7e-16 2e-16 \ 5: 8.0000e-02 7.9999e-02 1e-06 2e-16 1e-15 \ 6: 8.0000e-02 8.0000e-02 1e-08 8e-16 1e-15 \ Optimal solution found. \ x_test: [-0.40000001 4.20000012] \

Exercise 1.2: Find Q, c, A, B for input into cvxopt_solve_qp function (10 points)

```
In [48]: Q = None
c = None
A = None
B = None
### BEGIN SOLUTION
Q = np.eye(3);
Q[2,2] = 0;
c = np.zeros([3])
A = np.multiply(np.tile(-y,[1, 3]), np.concatenate([X, np.ones([2
0,1])],1))
B = -np.ones([20])
### END SOLUTION

x = cvxopt_solve_qp(Q, c, A, B)
```

	pcost	dcost	gap	pres	dres
0:	1.4721e-03	6.5053e+00	5e+01	2e+00	4e+02
1:	1.0012e-02	-4.7161e+00	1e+01	6e-01	1e+02
2:	2.6180e-02	-4.8172e+00	7e+00	2e-01	6e+01
3:	3.9767e-02	-4.5363e-01	5e-01	1e-02	2e+00
4:	3.5404e-02	1.8200e-02	2e-02	4e-15	2e-13
5:	3.1392e-02	3.0877e-02	5e-04	6e-15	2e-12
6:	3.1250e-02	3.1245e-02	5e-06	4e-15	6e-13
7:	3.1249e-02	3.1248e-02	5e-08	4e-15	2e-13

Optimal solution found.

```
In [49]: print('Q:\n', Q)
print('c:\n', c)
print('A:\n', A[7:13])
print('B:\n', B)
print('x:\n', x)
# Test function: Do not remove
assert Q.shape == (3, 3) and Q[2,2] == Q[0,1] and Q[2,0] == 0 and
Q[0,0] == Q[1,1] and Q[0,0] == 1, 'Q value is incorrect'
assert c.shape == 3 or c.shape == (3,) or c.shape == (3,1), 'Size
of c is incorrect'
assert np.all((c == 0)), 'c value is incorrect'
assert A.shape == (20,3), 'Size of A is incorrect'
assert np.max(A[:,2]) == 1 and np.min(A[:,2]) == -1, 'A value is
incorrect'
assert not np.array_equal(np.round(A[:,0:2],1), np.round(X,1)), '
A value is incorrect'
assert np.array_equal(np.round(x,1), np.round([0.16001143, 0.1920
7647, -38.32646165],1)), 'x value is incorrect'

print("success!")
# End Test function
```

```
Q:
[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 0.]]
c:
[0. 0. 0.]
A:
[[ 164.411   42.38    1.   ]
 [ 157.308   42.15    1.   ]
 [ 159.579   49.739    1.   ]
 [-168.524  -64.353   -1.   ]
 [-171.597  -61.793   -1.   ]
 [-179.469  -74.552   -1.   ]]
B:
[-1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1.
-1. -1.
-1. -1.]
x:
[ 0.16001143  0.19207647 -38.32646165]
success!
```

Expect Result (or look-alike): \ Q:\ [[1. 0. 0.] \ [0. 1. 0.] \ [0. 0. 0.]] \ c:\ [0. 0. 0.] \ A:\ [[164.411 42.38 1.] \ [157.308 42.15 1.] \ [159.579 49.739 1.] \ [-168.524 -64.353 -1.] \ [-171.597 -61.793 -1.] \ [-179.469 -74.552 -1.]] \ B:\ [-1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1.] \ x:\ [0.16001143 0.19207647 -38.32646165] \

Exercise 1.3: Use x from above to find w and b (5 points)

```
In [50]: w = None
b = None
### BEGIN SOLUTION
w = np.matrix([[x[0]], [x[1]]])
b = x[2]
scale = np.linalg.norm(w)
w = w / scale
b = b / scale
### END SOLUTION
```

```
In [51]: print('Optimal w: [%f %f] b: %f' % (w[0,0],w[1,0],b))

plot_mf(Xf,Xm)
plot_w(w,b)
plt.show()

# Test function: Do not remove
assert w.shape == 2 or w.shape == (2,) or w.shape == (2,1), 'Size
of w is incorrect'
assert w[0] > 0 and w[1] > 0 and w[0] <= 1 and w[1] <= 1, 'w valu
e is incorrect'
assert isinstance(b, (float, int)), 'Type of b is incorrect'
assert b < 0, 'b value is incorrect'


print("success!")
# End Test function
```

Optimal w: [0.640061 0.768324] b: -153.309495

<ipython-input-45-a17db04aa7fd>:14: MatplotlibDeprecationWarning:
Adding an axes using the same arguments as a previous axes current-
ly reuses the earlier instance. In a future version, a new inst-
ance will always be created and returned. Meanwhile, this warnin-
g can be suppressed, and the future behavior ensured, by passing
a unique label to each axes instance.


ylim = plt.axes().get_ylim()
<ipython-input-45-a17db04aa7fd>:15: MatplotlibDeprecationWarning:
Adding an axes using the same arguments as a previous axes current-
ly reuses the earlier instance. In a future version, a new inst-
ance will always be created and returned. Meanwhile, this warnin-
g can be suppressed, and the future behavior ensured, by passing
a unique label to each axes instance.

xlim = plt.axes().get_xlim()



success!

Expect result (Or look-alike): \ Optimal w: [0.640061 0.768324] b: -153.309495



```
In [52]: def predict_linear(X,w,b):
          s = X@w+b
          s[s >= 0] = 1
          s[s < 0] = -1
          return s
          y_pred = predict_linear(X,w,b)
          accuracy = np.sum(y_pred==y)/y.size
          print(accuracy)
```

1.0

Exercise 2 (in lab): non-separable data (5 points)

Take the example of the annulus from the logistic regression lab. Verify that cvxopt cannot find a decision boundary for this case. Show your results in your lab report.


Note: You don't need to separate data to train/test values

```
In [53]: # Generate data for class 1
mu_1 = np.array([1.0, 2.0])
sigma_1 = 1
num_sample = 100
cov_mat = np.matrix([[sigma_1,0],[0,sigma_1]])
X1 = np.random.multivariate_normal(mean= mu_1, cov=cov_mat, size
= num_sample)

# Generate data for class 2
angle = np.random.uniform(0, 2*np.pi, num_sample)
d = np.random.normal(np.square(3*sigma_1),np.square(.5*sigma_1),
num_sample)
X2 = np.array([X1[:,0] + d*np.cos(angle), X1[:,1] + d*np.sin(angl
e)]).T

# Combine X1 and X2 into single dataset
X_annulus = np.concatenate([X1, X2],axis = 0)
y_annulus = np.append(-np.ones(num_sample),np.ones(num_sampl
e))[:,np.newaxis]
```

```
In [54]: fig1 = plt.figure(figsize=(8,8))
ax = plt.axes()
plt.title('Sample data for classification problem')
plt.grid(axis='both', alpha=.25)
plt.plot(X1[:,0],X1[:,1],'b.', label = 'Class 1')
plt.plot(X2[:,0],X2[:,1],'g*', label = 'Class 2')
plt.legend(loc=2)
plt.axis('equal')
plt.show()
```



```
In [55]: # Try to use try/catch to get output
get_error = False
try:
    Q = None
    c = None
    A = None
    B = None
    x = None
    w = None
    b = None
    ### BEGIN SOLUTION
    Q = np.eye(3);
    Q[2,2] = 0;
    c = np.zeros([3])
    A = np.multiply(np.tile(-y_annulus,[1, 3]), np.concatenate([X
_annulus, np.ones([200,1]),1]))
    b = -np.ones([200])
    x = cvxopt_solve_qp(Q, c, A, b)
    w = np.matrix([[x[0]], [x[1]]])
    b = x[2]
    scale = np.linalg.norm(w)
    w = w / scale
    b = b / scale
    ### END SOLUTION
    output_str = 'Optimal w: [%f %f] b: %f' % (w[0,0],w[1,0],b)
    get_error = False
except Exception as e:
    output_str = e
    get_error = True
```

	pcost	dcost	gap	pres	dres
0:	8.3375e-05	1.9930e+02	2e+02	2e+00	6e-14
1:	7.6693e-09	4.6646e+02	5e+00	1e+00	2e-13
2:	7.2606e-13	4.1550e+04	5e+00	1e+00	6e-12
3:	7.2562e-17	3.6072e+08	4e+02	1e+00	2e-07
4:	7.2562e-21	3.1308e+14	4e+06	1e+00	3e-02
5:	7.2564e-25	2.7173e+22	3e+12	1e+00	3e+06
6:	2.1134e-26	2.3230e+32	7e+20	1e+00	3e+16
7:	2.2488e-26	2.7866e+40	8e+28	1e+00	1e+25
8:	2.9744e-26	1.0338e+47	3e+35	1e+00	2e+31
9:	6.1706e-26	4.9003e+54	1e+43	1e+00	2e+39
10:	5.3970e-26	1.4914e+62	4e+50	1e+00	9e+46
11:	5.1908e-26	6.2056e+69	2e+58	1e+00	3e+54
12:	4.6299e-26	1.5907e+77	5e+65	1e+00	6e+61
13:	5.5269e-26	4.1712e+84	1e+73	1e+00	3e+69
14:	4.2903e-26	1.9751e+92	6e+80	1e+00	4e+76
15:	5.7310e-26	2.2560e+99	6e+87	1e+00	7e+83
16:	4.8119e-26	5.2541e+106	1e+95	1e+00	3e+91
17:	5.4391e-26	1.4981e+114	4e+102	1e+00	6e+98
18:	4.2148e-26	6.9060e+121	2e+110	1e+00	1e+106
19:	5.6850e-26	1.0980e+129	3e+117	1e+00	5e+113
20:	4.5908e-26	4.0391e+136	1e+125	1e+00	2e+121
21:	5.5264e-26	1.2168e+144	3e+132	1e+00	9e+128
22:	4.9144e-26	4.0387e+151	1e+140	1e+00	1e+136
23:	5.7176e-26	1.2269e+159	3e+147	1e+00	7e+143
24:	5.6123e-26	2.5414e+166	7e+154	1e+00	3e+150
25:	5.6119e-26	7.2104e+173	2e+162	1e+00	inf
26:	4.6759e-26	1.8358e+181	5e+169	1e+00	inf
27:	5.4745e-26	5.0501e+188	1e+177	1e+00	inf
28:	4.3312e-26	2.3436e+196	7e+184	1e+00	inf
29:	5.7894e-26	2.6939e+203	8e+191	1e+00	inf
30:	4.8851e-26	5.9572e+210	2e+199	1e+00	inf
31:	5.3756e-26	1.7786e+218	5e+206	1e+00	inf
32:	4.2038e-26	7.8614e+225	2e+214	1e+00	inf
33:	5.6018e-26	1.5303e+233	4e+221	1e+00	inf
34:	4.1044e-26	6.7578e+240	2e+229	1e+00	inf
35:	5.5555e-26	1.3360e+248	4e+236	1e+00	inf
36:	4.1419e-26	5.9255e+255	2e+244	1e+00	inf
37:	5.5857e-26	1.0975e+263	3e+251	1e+00	inf
38:	4.1216e-26	4.7003e+270	1e+259	1e+00	inf
39:	5.5217e-26	1.0802e+278	3e+266	1e+00	inf
40:	4.4281e-26	4.9675e+285	1e+274	1e+00	inf

```
In [56]: print(output_str)

# Test function: Do not remove
assert Q.shape == (3, 3) and Q[2,2] == Q[0,1] and Q[2,0] == 0 and
Q[0,0] == Q[1,1] and Q[0,0] == 1, 'Q value is incorrect'
assert c.shape == 3 or c.shape == (3,) or c.shape == (3,1), 'Size
of c is incorrect'
assert np.all((c == 0)), 'c value is incorrect'
assert A.shape == (200,3), 'Size of A is incorrect'
assert str(output_str) == 'domain error' or "'NoneType' object is
not subscriptable" or get_error, 'Output incorrect'

print("success!")
# End Test function

domain error
success!
```

Expect result: Show something error calculation

Generalized Lagrangian optimization for SVMs

Now we consider the generalized Lagrangian for the SVM. This technique is suitable for solving problems of the form

$$\begin{aligned} \min_{\mathbf{w}} \quad & f(\mathbf{w}) \\ \text{subject to} \quad & g_i(\mathbf{w}) \leq 0, i \in 1..k \\ & h_i(\mathbf{w}) = 0, i \in 1..l \end{aligned}$$

The generalized Lagrangian is

$$\mathcal{L}(\mathbf{w}, \alpha, \beta) = f(\mathbf{w}) + \sum_{i=1}^k \alpha_i g_i(\mathbf{w}) + \sum_{i=1}^l \beta_i h_i(\mathbf{w}),$$

which has been cleverly arranged to be equal to $f(\mathbf{w})$ whenever \mathbf{w} satisfies the constraints and ∞ otherwise.

Primal and dual Lagrangian problems

The primal problem is to find

$$p^* = \min_{\mathbf{w}} \theta_{\mathcal{P}}(\mathbf{w}) = \min_{\mathbf{w}} \max_{\alpha, \beta, \alpha_i \geq 0} \mathcal{L}(\mathbf{w}, \alpha, \beta)$$

and the dual problem is to find

$$d^* = \max_{\alpha, \beta, \alpha_i \geq 0} \theta_{\mathcal{D}}(\alpha, \beta) = \max_{\alpha, \beta, \alpha_i \geq 0} \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}, \alpha, \beta).$$

If f is convex, the g_i 's are affine, the h_i 's are convex, and the g_i 's are strictly feasible, it turns out that the solutions to the primal and dual problem are the same, and the KKT conditions hold:

$$\begin{aligned} \frac{\partial}{\partial w_i} \mathcal{L}(\mathbf{w}^*, \alpha^*, \beta^*) &= 0, i \in 1..n \\ \frac{\partial}{\partial \beta_i} \mathcal{L}(\mathbf{w}^*, \alpha^*, \beta^*) &= 0, i \in 1..l \\ \alpha_i^* g_i(\mathbf{w}^*) &= 0, i \in 1..k \\ g_i(\mathbf{w}^*) &\leq 0, i \in 1..k \\ \alpha_i^* &\geq 0, i \in 1..k \end{aligned}$$

Solving the dual Lagrangian problem

The dual problem turns out to be easiest to solve.

We first solve for \mathbf{w} assuming fixed α and β (we don't have equality constraints though, so no need for β).

We need to rewrite the SVM constraints in the necessary form with $g_i(\mathbf{w}) = 0$. obtain, for the SVM, constraints

$$g_i(\mathbf{w}, b) = -y^{(i)}(\mathbf{w}^\top \mathbf{x}^{(i)} + b) + 1 \geq 0.$$

Using that definition of $g_i(\mathbf{w}, b)$, we obtain the Lagrangian

$$\mathcal{L}(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^m \alpha_i \left[y^{(i)}(\mathbf{w}^\top \mathbf{x}^{(i)} + b) - 1 \right]$$

Taking the gradient of \mathcal{L} with respect to \mathbf{w} and setting it to 0, we obtain

$$\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}, b, \alpha) = \mathbf{w} - \sum_{i=1}^m \alpha_i y^{(i)} \mathbf{x}^{(i)} = 0,$$

which gives us

$$\mathbf{w} = \sum_{i=1}^m \alpha_i y^{(i)} \mathbf{x}^{(i)}.$$

From $\frac{\partial \mathcal{L}}{\partial b} = 0$, we obtain

$$\sum_{i=1}^m \alpha_i y^{(i)} = 0,$$

which is interesting (think about what it means also considering that $\alpha_i > 0$ only for examples on the margin. Unfortunately it doesn't help us find b ! In any case, we plug this definition for the optimal \mathbf{w} into the original Lagrangian, to obtain

$$\mathcal{L}(\mathbf{w}, b, \alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j (\mathbf{x}^{(i)})^\top \mathbf{x}^{(j)} - b \sum_{i=1}^m \alpha_i y^{(i)}.$$

We already know that the last term is 0, so we get

$$\mathcal{L}(\mathbf{w}, b, \alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle \mathbf{x}^{(i)}, \mathbf{x}^{(j)} \rangle.$$

OK! We've eliminated \mathbf{w} and b from the optimization. Now we just need to maximize \mathcal{L} with respect to α .

This gives us the final (dual) optimization problem

$$\max_{\alpha} \quad W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle \mathbf{x}^{(i)}, \mathbf{x}^{(j)} \rangle$$

such that $\alpha_i \geq 0, i \in 1..m$

$$\sum_{i=1}^m \alpha_i y^{(i)} = 0$$

This turns out to be QP again!

Aside: once we solve for α , we obtain \mathbf{w} according to the equation above, then it turns out that the optimal b can be obtained as in the lecture notes.

QP solution to dual problem

We need to negate our objective function to turn the max (SVM formulation) into a min (QP formulation).

For the second term of $W(\alpha)$, first let \mathbf{K} be the kernel matrix with $K_{ij} = \langle \mathbf{x}^{(i)}, \mathbf{x}^{(j)} \rangle$. Then $\alpha^\top \text{diag}(\mathbf{y}) \mathbf{K} \text{diag}(\mathbf{y}) \alpha$ gives us the summation in the second term ($\text{diag}(\mathbf{y})$ is just the square diagonal matrix with \mathbf{y} as its diagonal).

The (negated) first term of $W(\alpha)$ can be written in QP form with $\mathbf{c} = [-1 \quad -1 \quad \dots]^\top$.

So that gives us our QP setup:

$$\begin{aligned} \mathbf{Q} &= \text{diag}(\mathbf{y}) \mathbf{K} \text{diag}(\mathbf{y}) & \mathbf{c} &= \begin{bmatrix} -1 \\ -1 \\ \vdots \end{bmatrix} \\ \mathbf{A} &= -\mathbf{I}_{m \times m} & \mathbf{b} &= \begin{bmatrix} 0 \\ 0 \\ \vdots \end{bmatrix} \\ \mathbf{E} &= \mathbf{y}^\top & \mathbf{d} &= [0]. \end{aligned}$$

OK, now the code:

```

In [94]: m = X.shape[0];
         n = X.shape[1];

# Transform data set so that each attribute has a
# mean of 0 and a standard deviation of 1

def preprocess(X):
    means = X.mean(0);
    scales = 1/np.std(X,0);
    Xh = np.concatenate([X.T,np.ones([1,20])],0);
    Tm = np.matrix(np.eye(3));
    Tm[0:2,2:3] = -X.mean(0).T;
    Ts = np.matrix(np.eye(3));
    Ts[0:2,0:2] = np.diagflat(scales);
    T = Ts*Tm;
    XX = (T * Xh);
    XX = XX[0:2,:].T;
    return XX, T;

# RBF/Gaussian kernel

def gauss_kernel(X):
    sigma = 0.2
    m = X.shape[0];
    K = np.matrix(np.zeros([m,m]));
    for i in range(0,m):
        for j in range(0,m):
            K[i,j] = (X[i,:] - X[j,:]).reshape(1,-1) @ (X[i,:] -
X[j,:]).reshape(-1,1)
    K = np.exp(-K/(2*sigma*sigma))
    return K;

def linear_kernel(X):
    m = X.shape[0];
    K = np.matrix(np.zeros([m,m]));
    for i in range(0,m):
        for j in range(0,m):
            K[i,j] = (X[i,:].reshape(1,-1)@X[j,:].reshape(-1,1))
    return K;

```

Exercise 3 (in lab): linearly separable data (15 points)

Take the example data from the exercise in lecture. Verify that you can use the dual optimization to find the decision boundary for such "easy" cases. Show your results in your lab report.

Exercise 3.1: Find Q, c, A, b, E, d for input into cvxopt_solve_qp function (10 points)

```
In [95]: Q = None
c = None
A = None
b = None
E = None
d = None

K = linear_kernel(X)
### BEGIN SOLUTION
Q = np.multiply(y * y.T, K)
c = -np.ones([m]);
A = -np.eye(m);
B = np.zeros([m]);
E = y.T;
d = np.zeros(1);
### END SOLUTION
alpha_star = cvxopt_solve_qp(Q, c, A, B, E, d)
```

	pcost	dcost	gap	pres	dres
0:	-2.7646e+00	-4.9725e+00	5e+01	6e+00	2e+00
1:	-6.4101e+00	-3.8299e+00	1e+01	2e+00	6e-01
2:	-5.0055e+00	-1.2719e+00	7e+00	8e-01	2e-01
3:	-5.1552e-02	-4.0648e-02	5e-01	3e-02	1e-02
4:	-1.8200e-02	-3.5404e-02	2e-02	2e-17	2e-13
5:	-3.0877e-02	-3.1392e-02	5e-04	8e-18	2e-13
6:	-3.1245e-02	-3.1250e-02	5e-06	4e-18	1e-13
7:	-3.1248e-02	-3.1249e-02	5e-08	3e-18	2e-13

Optimal solution found.

```
In [96]: print('Q rank: %d' % np.linalg.matrix_rank(Q))
print("Optimal alpha:\n", alpha_star)

# Test function: Do not remove
assert Q.shape == (20, 20), 'Size of Q is incorrect'
assert np.linalg.matrix_rank(Q) == 2, 'Q rank is incorrect'
assert np.all((c == -1)), 'c value is incorrect'
assert A.shape == (20,20), 'Size of A is incorrect'
assert np.all((B == 0)), 'b value is incorrect'
assert np.array_equal(np.round(E,1), np.round(y.T,1)), 'E value is incorrect'
assert d.shape == (1,) or d.shape == 1 or d.shape == (1,1), 'Size of d is incorrect'
assert np.all((d == 0)), 'd value is incorrect'
assert alpha_star.shape == (20,) or alpha_star.shape == 20 or alpha_star.shape == (20,1), 'Size of alpha_star is incorrect'

print("success!")
# End Test function
```

```
Q rank: 2
Optimal alpha:
[3.12484796e-02 1.13821985e-09 7.68004003e-10 6.22346942e-10
 6.43114906e-10 8.93974184e-10 6.60695980e-10 4.31409827e-10
 6.68557407e-10 1.19689394e-09 1.56332821e-02 1.56151999e-02
 5.00606721e-10 6.71834910e-10 4.89168003e-10 4.93113729e-10
 4.90935291e-10 4.85894509e-10 9.54348935e-10 4.42078646e-10]
success!
```

Expect Result (or look a like): \ Q rank: 2\ Optimal alpha:\ [3.12484796e-02 1.13821985e-09
 7.68004003e-10 6.22346942e-10\ 6.43114906e-10 8.93974184e-10 6.60695980e-10 4.31409827e-10\
 6.68557407e-10 1.19689394e-09 1.56332821e-02 1.56151999e-02\ 5.00606721e-10 6.71834910e-10
 4.89168003e-10 4.93113729e-10\ 4.90935291e-10 4.85894509e-10 9.54348935e-10 4.42078646e-10]

Exercise 3.2: write get_wb function (5 points)

```
In [97]: def get_wb(X, y, alpha, K):
# Find the support vectors
S = alpha > 1e-6
XS = None
yS = None
alphaS = None
alphaSyS = None
w = None
# Find b
KS = None
NS = None
b = None
# Normalize w,b
scalef = None
w = None
b = None

### BEGIN SOLUTION
S = alpha > 1e-6
XS = X[S,:]
yS = y[S]
alphaS = alpha[S]
alphaSyS = np.tile(np.multiply(yS.T, alphaS).T, n)
w = sum(np.multiply(alphaSyS, XS)).T
# Find b
KS = K[S,:][:,S]
NS = yS.shape[0]
b = (np.sum(yS) - np.sum(np.multiply(alphaS,yS.T)*KS))/NS
# Normalize w,b
scalef = np.linalg.norm(w)
w = w / scalef
b = b / scalef
### END SOLUTION
return w,b
```

```
In [98]: # Test function: Do not remove
w,b = get_wb(X, y, alpha_star, K)

print("Optimal w: [%f,%f] b: %f" % (w[0],w[1],b))
plot_mf(Xf,Xm)

plot_w(w,b)
plt.show()

print("success!")
# End test function
```


Optimal w: [0.640062,0.768324] b: -153.309583

<ipython-input-45-a17db04aa7fd>:14: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ylim = plt.axes().get_ylim()
```


<ipython-input-45-a17db04aa7fd>:15: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
xlim = plt.axes().get_xlim()
```



success!

Expect Result (Or look-alike): \ Optimal w : [0.640062,0.768324] b : -153.309583



```
In [99]: y_pred = predict_linear(X,w,b)
accuracy = np.sum(y_pred==y)/y.size
print(accuracy)
```

1.0

Exercise 4 (in lab): non-separable data, linear kernel (15 points)

Again, take the example of the annulus from the logistic regression lab. Verify that the dual optimization with the linear kernel still cannot find a decision boundary for this case. Show your results in your lab report.

```
In [118]: (m, n) = X_annulus.shape
print(X_annulus.shape)
K_annulus = linear_kernel(X_annulus);

Q_annulus = None
c = None
A = None
B = None
E = None
d = None

### BEGIN SOLUTION
Q_annulus = np.multiply(y_annulus * y_annulus.T, K_annulus)

c = -np.ones([m]);
A = -np.eye(m);
B = np.zeros([m]);
E = y_annulus.T;
d = np.zeros(1);
### END SOLUTION

(200, 2)
```

```
In [119]: alpha_star_annulus = cvxopt_solve_qp(Q_annulus, c, A, B, E, d)

print('Q rank: %d' % np.linalg.matrix_rank(Q_annulus))
print("Optimal alpha:", alpha_star_annulus)
# Test function: Do not remove
assert np.linalg.matrix_rank(Q_annulus) > 2, "Q rank is incorrect"
assert alpha_star_annulus is None, "alpha_star_annulus cannot be calculated."

print("success!")
# End test function
```

	pcost	dcost	gap	pres	dres
0:	-1.9930e+02	-4.2459e+02	2e+02	8e-15	2e+00
1:	-4.6646e+02	-4.7177e+02	5e+00	6e-14	1e+00
2:	-4.1550e+04	-4.1555e+04	5e+00	4e-11	1e+00
3:	-3.5661e+08	-3.5661e+08	9e+02	5e-07	1e+00
4:	-4.3418e+08	-4.3418e+08	1e+03	3e-07	1e+00

Terminated (singular KKT matrix).
Q rank: 2
Optimal alpha: None
success!

```
In [120]: get_error = False
try:
    w,b = get_wb(X_annulus, y_annulus, alpha_star_annulus, K_annulus)

    output_str = "Optimal w: [%f,%f] b: %f" % (w[0],w[1],b)
    plot_mf(Xf,Xm)
    plot_w(w,b)
    get_error = False
except Exception as e:
    output_str = str(e)
    get_error = True

print(output_str)
# Test function: Do not remove
assert str(output_str) == 'domain error' or "'NoneType' object is
not subscriptable" or get_error, 'Output incorrect'

print("success!")
# End Test function

'>' not supported between instances of 'NoneType' and 'float'
success!
```

Expect Result: Any error

Exercise 5 (in lab): "easy" non-separable data, Gaussian (RBF) kernel with non-overlapping data (10 points)

Now use the Gaussian (radial basis function) kernel instead of the linear kernel implemented in the code above and verify that you can correctly solve the problem.

```
In [129]: (m, n) = X_annulus.shape
K_annulus = None
Q_annulus = None
c = None
A = None
B = None
E = None
d = None

### BEGIN SOLUTION
K_annulus = gauss_kernel(X_annulus)

Q_annulus = np.multiply(y_annulus * y_annulus.T, K_annulus)

c = -np.ones([m]);
A = -np.eye(m);
B = np.zeros([m]);
E = y_annulus.T;
d = np.zeros(1);
### END SOLUTION
```



```
In [130]: alpha_star_annulus = cvxopt_solve_qp(Q_annulus, c, A, B, E, d)
w,b = get_wb(X_annulus, y_annulus, alpha_star_annulus, K_annulus)
w = w.reshape(-1,1)

print('Q rank: %d' % np.linalg.matrix_rank(Q))
print("Optimal alpha:")
print(alpha_star_annulus[:5])
print(w.shape)
print("Optimal w: [%f,%f] b: %f" % (w[0],w[1],b))

# Test function: Do not remove
assert np.linalg.matrix_rank(Q_annulus) > 2, "Q rank is incorrect"
assert alpha_star_annulus is not None, "alpha_star_annulus cannot be calculated."
assert w.shape == (2,1), 'Size of w is incorrect'
assert np.all(w <= 1) and np.all(w >= -1), 'w value is incorrect'

print("success!")
# End test function
```

```
      pcost      dcost      gap      pres      dres
0: -5.0795e+01 -1.3928e+02 9e+01 1e-15 2e+00
1: -5.9073e+01 -6.6466e+01 7e+00 2e-14 3e-01
2: -6.2076e+01 -6.3563e+01 1e+00 1e-14 4e-02
3: -6.2184e+01 -6.2423e+01 2e-01 8e-15 5e-03
4: -6.2205e+01 -6.2232e+01 3e-02 3e-14 5e-04
5: -6.2209e+01 -6.2209e+01 7e-04 4e-14 6e-06
6: -6.2209e+01 -6.2209e+01 1e-05 1e-14 7e-08
Optimal solution found.
Q rank: 2
Optimal alpha:
[5.77224260e-01 1.31648362e+00 1.17475881e+00 4.38419640e-01
 8.16578932e-07]
(2, 1)
Optimal w: [-0.974636,0.223795] b: 0.008878
success!
```

Expect result (or look-alike): \ pcost dcost gap pres dres \ 0: -5.0795e+01 -1.3928e+02 9e+01 1e-15 2e+00 \ 1: -5.9073e+01 -6.6466e+01 7e+00 2e-14 3e-01 \ 2: -6.2076e+01 -6.3563e+01 1e+00 1e-14 4e-02 \ 3: -6.2184e+01 -6.2423e+01 2e-01 8e-15 5e-03 \ 4: -6.2205e+01 -6.2232e+01 3e-02 3e-14 5e-04 \ 5: -6.2209e+01 -6.2209e+01 7e-04 4e-14 6e-06 \ 6: -6.2209e+01 -6.2209e+01 1e-05 1e-14 7e-08 \ Optimal solution found. \ Q rank: 2 \ Optimal alpha: \ [5.77224260e-01 1.31648362e+00 1.17475881e+00 4.38419640e-01 8.16578932e-07] \ (2, 1) \ Optimal w: [-0.974636,0.223795] b: 0.008878

```
In [136]: def predict(x,X,y,alpha):  
    s = []  
    sigma = 0.2  
    for j in range(x.shape[0]):  
        ss = 0  
        for i in range(X.shape[0]):  
            ss += alpha[i]*y[i]*np.exp(-(X[i]-x[j])@(X[i]-x  
[j]))/(2*sigma*sigma))  
        s.append(ss)  
    s = np.array(s)  
    s[s >= 0] = 1  
    s[s < 0] = -1  
    return s  
  
y_pred = predict(X_annulus,X_annulus,y_annulus,alpha_star_annulu  
s)  
np.sum(y_annulus == y_pred)/y_annulus.size
```

Out[136]: 1.0

See your graph of classification

```
In [137]: x_series = np.linspace(-15,15,100)
y_series = np.linspace(-15,15,100)

x_mesh,y_mesh = np.meshgrid(x_series,y_series)


x_mesh = x_mesh.reshape(-1,1)
y_mesh = y_mesh.reshape(-1,1)


mesh = np.append(x_mesh,y_mesh,axis=1)
y_pred = predict(mesh,X_annulus,y_annulus,alpha_star_annulus)

x_mesh = x_mesh.reshape(100,100)
y_mesh = y_mesh.reshape(100,100)
y_pred = y_pred.reshape(100,100)

#plt.scatter(mesh[:,0],mesh[:,1],c=y_pred)
plot_mf(X1,X2)
plt.pcolormesh(x_mesh,y_mesh,y_pred,cmap='viridis',shading='auto',alpha=0.1)
```

Out[137]: <matplotlib.collections.QuadMesh at 0x7fb4664480d0>



Expect Result:**Exercise 6 (take home): more difficult non-separable data**

Now find or generate a dataset in which the decision boundary is nonlinear AND the data overlap along that nonlinear boundary. Show that the result.

In []:

In []:

In []:

In []: